



P3 AI-DLC → P3 ADLC

The P3 Methodology

From AI-assisted development to governed Agentic Development Lifecycle.

P3 SOLUTIONS GROUP

Integrating People, Practices, and Principles

P3 AI-DLC → P3 ADLC

A maturity path for moving from individual AI assistance to governed, measurable agentic software delivery.

AI-DLC defines how AI participates across the software lifecycle. P3 ADLC is the mature operating model in which agents can orchestrate bounded portions of that lifecycle under engineered controls and explicit authority.

Level	Operating model	Primary shift
L0	Traditional SDLC	Human-driven delivery
L1	AI-Assisted	Individual productivity
L2	Standardized AI Development	Shared practices and controls
L3	P3 AI-DLC	AI participates across the lifecycle
L4	Agentic AI-DLC	Agents orchestrate bounded work
L5	P3 ADLC	Governed agentic delivery system

The P3 lifecycle

Intent → Spec → Readiness → Plan → Build → Verify → Review → Release → Observe → Learn

The lifecycle is intentionally broader than coding. P3 focuses on the entire path from business intent to production evidence, because local coding acceleration can simply move bottlenecks downstream if the surrounding delivery system does not change.

The Control System Around Probabilistic AI

The Engineering Harness is the deterministic context and control environment in which AI and agents operate.

CONTEXT

Specs and product intent

- Architecture decisions
- Repository memory
- Standards and conventions
- Runbooks and operating knowledge

CAPABILITIES

Commands and skills

- MCP/tool access
- Worktrees/containers
- CI/CD integration
- Agent roles

CONSTRAINTS

Decision rights

- Security boundaries
- Branch protections
- Dependency policies
- Production authority

FEEDBACK

Lint/typecheck/tests

- Quality gates
- Build evidence
- SRE/production telemetry
- Measured delivery outcomes

When an agent repeatedly fails, P3 does not default to 'write a better prompt.' We ask what capability, context, constraint, tool, or feedback loop is missing from the harness.

From Human Gates to Earned Autonomy

Early AI-DLC contains more explicit human gates. That is deliberate. Each gate identifies a decision involving judgment, risk, or authority. Over time, evidence determines which gates can become policy-driven and which should remain human.

Example progression

- **Stage 1:** Human initiates a bounded spec; AI performs readiness analysis; human approves architecture; AI implements; CI verifies; human approves release.
- **Stage 2:** Product initiates from Slack; an intake agent creates/refines the spec; engineering agents proceed within codified authority; humans intervene only at material decision boundaries.
- **Stage 3:** Agents orchestrate approved lifecycle segments, remediate deterministic failures, collect evidence, and stop only at defined authority gates.

P3 measurement loop

PAIN → MEASURE → BUILD → PROVE → PACKAGE → SELL

For engineering methodology, the learning loop is: **Observed → Hypothesis → Tested → Validated → P3 Standard**. We do not declare a practice 'best' because it sounds sophisticated. We run it, measure it, and codify what survives contact with real delivery.

The Abstraction Layer Moves Up

Developer / Engineer

Less time spent on repetitive production of code. More time defining intent, constraints, architecture, verification, and the harness that enables AI to produce safely.

Platform / DevOps

CI/CD becomes the independent enforcement layer around AI-generated work. Deterministic verification becomes more valuable as generation becomes more probabilistic.

Engineering Leader

Less emphasis on policing individual AI usage. More emphasis on operating model, decision rights, delivery economics, quality gates, and evidence.

Product

Product intent becomes increasingly machine-actionable: clear outcomes, scope, acceptance criteria, constraints, and unresolved decisions become the contract that initiates the lifecycle.

P3 AI-DLC is not a tool rollout. P3 ADLC is not unattended autonomy. Both are engineering-system transformations.